

Local Vision OCR Benchmark

3 receipts × 3 runs · NVIDIA RTX A2000 8GB Laptop GPU, 8192 MiB, 610.43.03 · Flash Attention false · f16 KV

MODELS

15

TRIALS

135

TOTAL TIME

39.0 min

RECOMMENDATION

Qwen 3.5 4B

#	Model	Exact	Stable	Warm	Cold	Size	Err
1	Qwen 3.5 4B	9/9	3/3	0.46s	6.23s	3.4G	0
2	Ministral 3 3B	6/9	2/3	0.32s	5.82s	3.0G	0
3	Qwen 3.5 2B	6/9	2/3	0.40s	5.18s	2.7G	0
4	Gemma 3 4B	6/9	2/3	0.53s	8.73s	3.3G	0
5	MiniCPM-V 2.6 8B	6/9	2/3	0.68s	31.22s	5.5G	0
6	MiniCPM-V 4.5 8B	6/9	2/3	0.70s	6.53s	6.1G	0
7	Qwen 3.5 0.8B	3/9	1/3	0.36s	13.61s	1.0G	0
8	MiniCPM-V 4.6 1B	3/9	1/3	0.88s	52.65s	1.6G	0
9	Moondream 2 1.8B	0/9	0/3	0.14s	2.80s	1.7G	0
10	GLM-OCR	0/9	0/3	0.75s	29.90s	2.2G	0
11	Qwen3-VL 2B	0/9	0/3	0.95s	5.22s	1.9G	0
12	Qwen3-VL 4B	0/9	0/3	1.69s	7.02s	3.3G	0
13	Granite 3.2 Vision 2B	0/9	0/3	—	—	2.4G	9
14	DeepSeek-OCR 3B	0/9	0/3	—	—	6.7G	9
15	Llama 3.2 Vision 11B	0/9	0/3	—	—	7.8G	9

Key finding

Qwen 3.5 4B is the recommended smallest near-fastest perfect model.

Smoke test only: three annotated German receipts; errors count as incorrect. Warm median excludes the first model request.